MATH 510 Course Project: Rotten or fresh?
Professor Kevin Huggins
based on a project created by Dr. Mine Çetinkaya-Rundel

You and your teammates work for Paramount Pictures. Your boss has just acquired data about audience movie scores along with a variety of explanatory variables. He is interested in learning what attributes make a movie popular and in predicting future audience scores. She is also interested in learning something new about movies. He wants your team to figure it all out.

The data can be loaded directly in RStudio using the following command:

load(url("https://stat.duke.edu/~mc301/data/movies.Rdata"))

Tasks

EDA: Do some exploratory data analysis to tell an "interesting" story about movies. Instead of limiting yourself to relationships between just two variables, broaden the scope of your analysis and employ creative approaches that evaluate relationships between two variables while controlling for another one.

Inference: Come up with a research question that can be answered with a hypothesis test or a confidence interval, e.g. "Is there a difference in mean audience scores between genres?" or "What is the average difference in audience scores between movies that do and do not feature without oscar winner actors?" This question could be used to shed some light on your choice of the "best" linear model. Carry out the appropriate inference task to answer your question.

Modeling: Develop the "best" multiple linear regression model to explain audience movie scores.

Prediction: Pick a movie from 2015 (a new movie that is not in the sample) and predict the audience score of that movie using the predict function in R. Also quantify the uncertainty around this prediction.

Data

The data set is comprised of 456 movies randomly sampled between 1974 and 2015.

Response: audience scores (average audience scores on a 0-100 scale).

Explanatory: there are several explanatory variables, such as genre, mpaa_rating, critics_score, etc. Details are given below.

Getting started

Create a new folder called "Final Project" in your Files tab, once created click on this folder, and then Click on *More -> Set As Working Directory*. This will ensure that your work for Final Project is saved in a separate folder from your other coursework.

Deliverables

A. Report

- ○ Due by midnight 13 April 2016
- ○ 3 pages max @ 12 pt font size, arial or times new roman font type
- ○ Your report should be organized with the following parts included and clearly labeled:
    1. **Introduction:** a summary of the data set and your goal.
    2. **EDA:** any univariate or bivariate summaries worth reporting.
    3. **Inference:** Answer the research question you have posed using a hypothesis test or a confidence interval.
    4. **The "Best" Model:**

        What is the "best" linear model for predicting the response variable? You do not need to explain every step you took to arrive at this model, but should give some indication of why you chose the model you did. If you tried a few different models, how did you settle on one?

        - How well does your model do? What is the percent variation explained?
        - What does your model tell you about relationships between your explanatory variables and your response variable?
        - What conditions do you need for your analysis to hold? What are the implications if some of those conditions are violated.
    5. **Prediction:**

Using your best model and a movie from late 2015 or 2016 of your choosing, predict its audience score. (You will need to find the relevant information about the movie online.)

Include a description of the uncertainty of you prediction.

6. *Conclusion:*
   ● What is the bottom line from your analysis?
   ● How well can you predict audience scores?
   ● What are the caveats to your analysis?
   ● Does this data set lack information that you would have liked to use?

C. Code
   ○ Due by midnight 13 April 2016
   ○ Additional details provided below

D. Presentation
   ○ Slides due by midnight 13 April 2016
   ○ To be given in the dates of 14-28 April, excluding Sat and Sun
   ○ 15 minutes max
   ○ Live synchronous delivery on Adobe connect with ALL team members present
   ○ Scheduling instructions will be provided later

Tips for your report and presentation

This project is an opportunity to apply what you have learned about descriptive statistics, graphical methods, correlation and regression, and hypothesis testing and confidence intervals.

The goal is not to do an exhaustive data analysis i.e., do not calculate every statistic and procedure you have learned for every variable, but rather to show that you are proficient at using R at a basic level and that you are proficient at interpreting and presenting the results.

You might consider critiquing your own method, such as issues pertaining to the reliability of the data and the appropriateness of the statistical analysis you used within the context of this specific data set.

Grading

Grading of the project will take into account:

Correctness: Are the procedures and explanations correct?
Presentation: What was the quality of the presentation and poster?
Content/Critical thought: Did your think carefully about the problem?

Your grade will be roughly based on the following components:

40% report
30% presentation
20% code
10% team peer evaluations

Team peer evaluation: You will be asked to fill out a survey where you rate the contribution of each team member. Filling out the survey is a prerequisite for getting credit on the team member evaluation. For grades less than 3.0, please provide some explanation. If any individual gets an average peer grade less than 2.0, this person will receive half the grade of the rest of the group.

Timeline and deadlines

Date	Task	Notes
19 Feb	Send team name and members to Professor Kevin	Each team is made up of 4 people. You can choose anyone from your current class
13 Apr	Presentation slides submitted to Moodle	
13 Apr	Project report submitted to Moodle	
13 Apr	Project code submitted to Moodle	
14-28 Apr	Project presentations	

Honor code

You may not discuss this project in any way with anyone outside your team, besides the professor. Failure to abide by this policy will result in a 0 for all teams involved.

Submitting code

Each team needs to submit code that can successfully be pasted into the following R code block. I should be able to run the code you submit and then print the summary of the multiple linear regression using the variable named mlr.

BEGIN YOUR CODE

# load data
load(url("https://stat.duke.edu/~mc301/data/movies.Rdata"))

# load the inference function if needed
load(url("https://stat.duke.edu/~mc301/R/fun/inference.RData"))

... anything you do prior to running the regression ...

mlr = lm(..., data = movies)

summary(mlr)

... anything you do after running the regression ...

END YOUR CODE

To submit your team's code, create an R Script called "teamname.R" where "teamname" is replaced with your team's name.

Your code must run with no errors on my computer with only the data file as the input. This means you must make sure any variables you define, and model you fit, etc. are all included in the code. Instructions for checking the reproducibility of your code are given below.

Note the variables that you are not allowed to use for prediction at the bottom of this document. Check that your analysis is fully reproducible with the code you submit

You can make sure that this script is working properly by doing the following.

First, save your workspace by running the following command in the console.

```
save.image("prj_workspace.RData")
```

(Alternatively, you can go to *Session -> Save Workspace As…* and then give it a name, e.g. prj_workspace.)

Then clear your workspace. Attention: The following command will delete every variable in your workspace, which is why it is important to do Step 1 first.

```
rm(list=ls())
```

(Alternatively you can click on the Clear button in your environment – the one with the broom icon.)

Load the dataset, run all of the code in your R script, and finally print out a summary of your model.

```
source("teamname.R")
summary(mlr)
```

If you get an error, go through the script "teamname.R" line by line (executing each line one at a time) to see where the error occurs. Note that one potential source of error might be that you named your final model something other than mlr.

To restore your workspace, do

```
load("prj_workspace.RData")
```

(Alternatively, you can go to Session -> Load Workspace… and choose the workspace you want to load, e.g. prj_workspace.)

This is the code I will be using to make sure you have submitted working code. Note that the file "movies.RData" needs to be in the same folder as "teamname.R" for the code above to work. Data details

The data frame has 25 attributes, which may be broken into three categories.

Title and response variable:

title: Title of movie

audience_score: Audience score on Rotten Tomatoes (response variable)

Explanatory variables you may use in the model.

type: Type of movie (Documentary, Feature Film, TV Movie)

genre: Genre of movie (Action & Adventure, Comedy, Documentary, Drama, Horror, Mystery & Suspense, Other)

runtime: Runtime of movie (in minutes)

year: Year the movie is released

mpaa_rating: MPAA rating of the movie (G, PG, PG-13, R, Unrated)

studio: Studio that produced the movie

imdb_num_votes: Number of votes on IMDB

critics_score: Critics score on Rotten Tomatoes

critics_rating: Categorical variable for critics rating on Rotten Tomatoes (Certified Fresh, Fresh, Rotten)

best_pic_nom: Whether or not the movie was nominated for a best picture Oscar (no, yes)

best_pic_win: Whether or not the movie won a best picture Oscar (no, yes)

best_actor_win: Whether or not one of the main actors in the movie ever won an Oscar (no, yes) – note that this is not necessarily whether the actor won an Oscar for their role in the given movie

best_actress win: Whether or not one of the main actresses in the movie ever won an Oscar (no, yes) – not that this is not necessarily whether the actresses won an Oscar for their role in the given movie

best_dir_win: Whether or not the director of the movie ever won an Oscar (no, yes) – not that this is not necessarily whether the director won an Oscar for the given movie

top200_box: Whether or not the movie is in the Top 200 Box Office list on BoxOfficeMojo (no, yes)

audience_rating: Categorical variable for audience rating on Rotten Tomatoes (Spilled, Upright)

director: Director of the movie

actor1-actor5: List of first 5 main actors in the movie (abridged cast), this information was used to determine whether the movie casts an actor or actress who won a best actor or actress Oscar

imdb_url: Link to IMDB page for the movie

rt_url: Link to Rotten Tomatoes page for the movie

imdb_id: IMDB ID of the movie

Some of these variables won't make sense to include in the model, and it's up to you to decide which ones might not make sense to include. And audience_rating should definitely not be included – that would be cheating!